ADSP 31014 Statistical Models for Data Science

Course Project Part 2 - Hira Stanley

Business Problem

Political Analysts would like to understand what factors drove the results of the 2020 U.S. Presidential election (specifically the percentage of votes for Joe Biden in each county), such as demographics, prior election results, and the impact of COVID.

Model Selection

The nature of this data set allows for either logistic or linear regression models to be performed. With logistic, we could predict a binary outcome of the 2020 election. However, because we also have access to percentage of vote share for the 2020 winner in each county, it would be more helpful in the context of the problem to do a linear regression so we can see how the other factors affect the vote share rather than only the final outcome.

Linear Regression Model

We build a linear regression model using a cleaned dataset of 3,021 rows. The predicted variable is percentage20_Joe_Biden. The model formula is:

```
percentage20_Joe_Biden ~ factor(state)*NonWhite + case_pct*Black + turnout_difference + Income + Men_pct + Transit
```

This model has 107 parameters with a R^2 of 0.76. See Appendix 1 for model summary.

Multicollinearity and Feature Engineering

We divide our 54 potentially explanatory variables into five sections to investigate and omit any columns that represent multicollinearity with each other or our predicted variable:

- Election results: The glaring multicollinearity is the 2016 election results with 0.98 correlation to our predicted variable. The simplest model would just be a linear relationship between percentage16_Hillary_Clinton and our predicted variable, which offers a r-squared of 0.95, but does not help us understand any other factors. We also omit total_votes2020, as it is multicollinear with total_votes16 (0.96) and turnout_difference (0.78).
- Demographic: TotalPop and VotingAgeCitizen have perfect correlation of 1, we remove TotalPop because in this context, the vote share should represent only those who are eligible to vote. There are 5 columns representing racial demographics, and none stand out with high multicollinearity that would inhibit explainability. We feature engineer a column called NonWhite, which sums all racial categories that are not White to represent a percentage of the population. A heatmap of correlation showed other potential multicollinearities, that allow us to reduce the 6 income-related columns to just income.
- Commute: Transit would offer some explainability without high correlation.
- Employment: Knowing the context of the dataset, we just keep Unemployment.
- Others: We previously converted COVID cases and deaths into a percentage of the population in that county (case_pct and death_pct) to not be multicollinear with TotalPop or VotingAgeCitizen. We also previously converted gender in a similar way to just represent the male population as a percentage (Men_pct).

Correlation Matrix Heatmap of Demographic Data (remainder in Appendix 2)

Interaction Terms

Categorical field state interacts with percentage of the population being NonWhite.

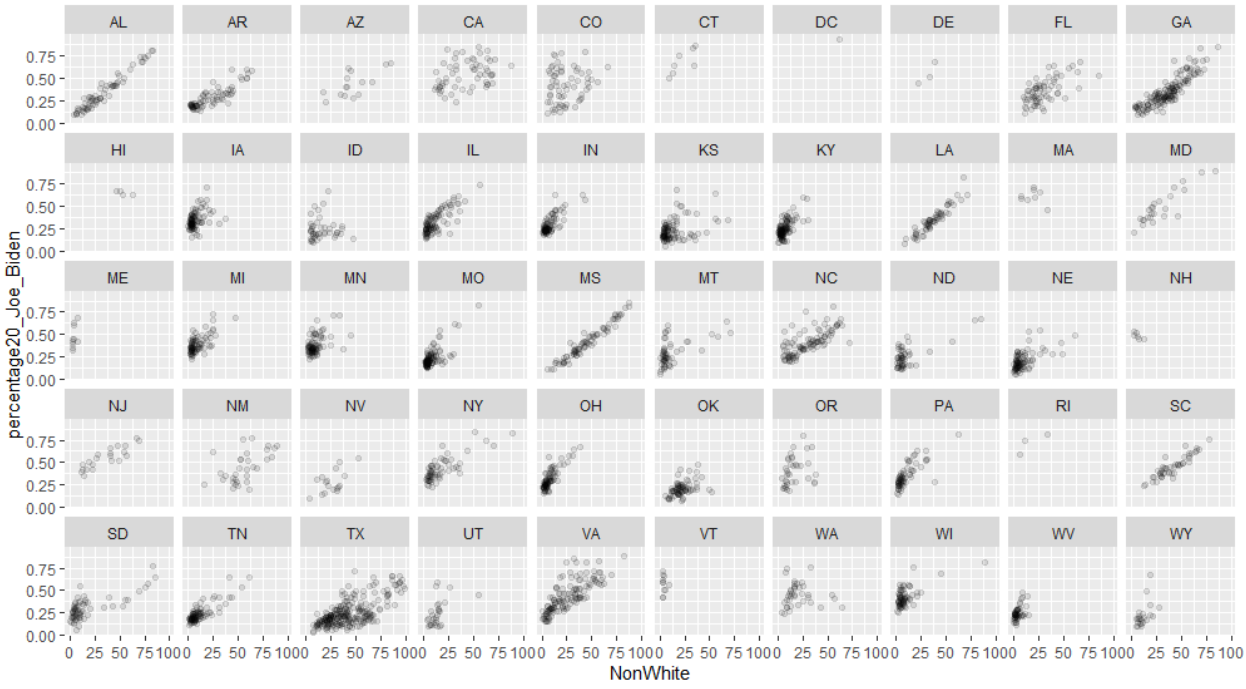

Useful and Less Useful Terms

Interaction terms:

- State interacting with the NonWhite percentage of the population significantly impacts the response variable. The p-values that are less than 0.05 in16 states indicate that the effect of NonWhite varies by state, with both positive and negative relationships observed.
- COVID cases (case_pct) when interacting with race had p-value 9.73e-05 and the final model has adjusted R^2 of 0.76. When input into the model separately without interaction, the adjusted R^2 is 0.75. This means the interaction term is more useful in providing a nuanced understanding of the data.

Less useful terms that we omitted either due to high p-values, or that their removal did not change the model's explanatory power (adjusted R^2).

- Unemployment had p-value 0.04
- Lat had p-value 0.26

Residuals

Plot of the response variable vs. residuals shows the model is a good fit, with the residuals randomly scattered around the horizontal line.

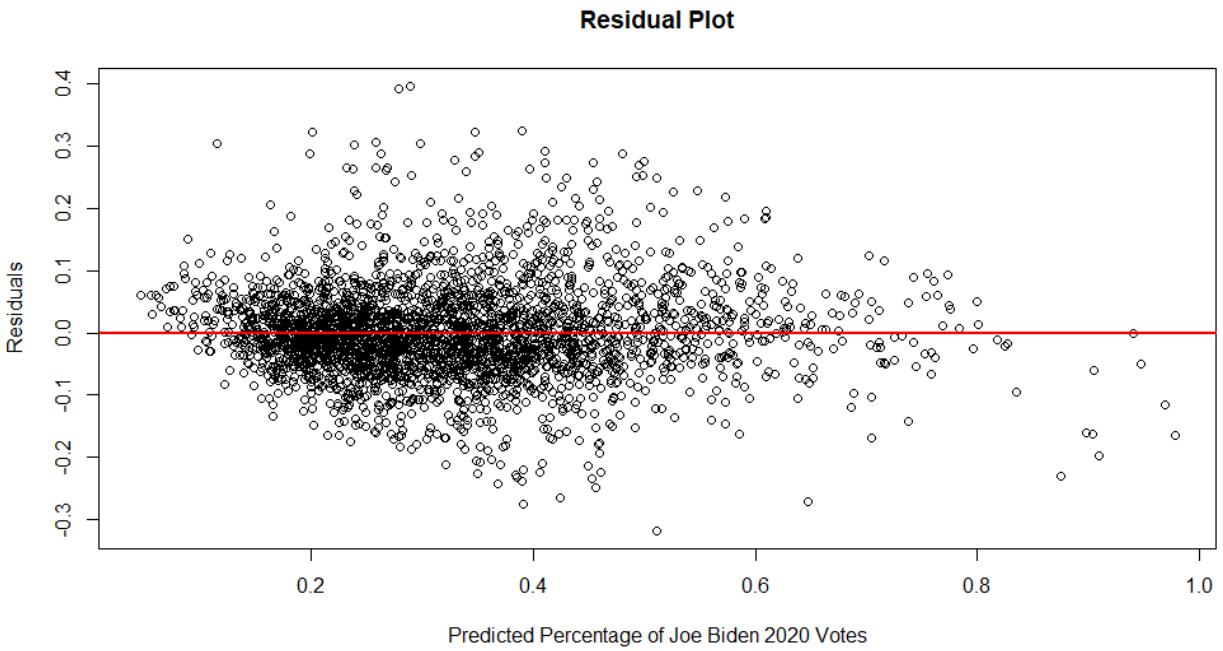

Model Trade-offs

The model where state interacted with NonWhite and case_pct had a slightly higher R^2 at 0.77, but the three-way interaction created 204 parameters. In the trade-off between model complexity and achieving a better fit, we choose the less complex one with 108 parameters.

Heteroscedasticity is present in multiple explanatory variables. When we take the log of the response variable, it reduces the R^2 down to 0.70. Despite the violation of the homoscedasticity assumption, we prioritize prediction accuracy in this trade-off. More detail in Appendices 3 and 4.

Appendix 1: Final Model Summary

```
Call:
lm(formula = percentage20_Joe_Biden ~ factor(state) * NonWhite +
    case_pct * Black + turnout_difference + Income + Men_pct +
    Transit, data = data)

Residuals:
     Min       1Q   Median       3Q      Max
-0.31771 -0.04380 -0.00571  0.03599  0.39612

Coefficients: (1 not defined because of singularities)
                       Estimate Std. Error t value Pr(>|t|)
(Intercept)            4.060e-01  3.806e-02  10.667  < 2e-16 ***
factor(state)AR        1.146e-01  2.289e-02   5.004 5.95e-07 ***
factor(state)AZ        4.532e-02  5.893e-02   0.769 0.441999
factor(state)CA        3.280e-01  3.045e-02  10.773  < 2e-16 ***
factor(state)CO        2.267e-01  2.641e-02   8.583  < 2e-16 ***
factor(state)CT        3.261e-01  7.152e-02   4.559 5.35e-06 ***
factor(state)DC       -2.071e-02  8.452e-02  -0.245 0.806447
factor(state)DE        9.797e-02  2.141e-01   0.458 0.647212
factor(state)FL        1.270e-01  2.758e-02   4.604 4.33e-06 ***
factor(state)GA       -2.923e-03  2.259e-02  -0.129 0.897044
factor(state)HI        1.037e+00  3.439e-01   3.016 0.002587 **
factor(state)IA        2.062e-01  2.146e-02   9.607  < 2e-16 ***
factor(state)ID        1.591e-01  2.692e-02   5.910 3.83e-09 ***
factor(state)IL        1.480e-01  2.123e-02   6.972 3.85e-12 ***
factor(state)IN        1.264e-01  2.155e-02   5.863 5.06e-09 ***
factor(state)KS        9.636e-02  2.109e-02   4.570 5.09e-06 ***
factor(state)KY        1.167e-01  2.125e-02   5.493 4.29e-08 ***
factor(state)LA       -7.022e-02  3.358e-02  -2.091 0.036632 *
factor(state)MA        4.392e-01  6.809e-02   6.450 1.30e-10 ***
factor(state)MD        1.480e-01  3.447e-02   4.294 1.82e-05 ***
factor(state)ME        2.211e-01  6.196e-02   3.568 0.000365 ***
factor(state)MI        2.420e-01  2.252e-02  10.744  < 2e-16 ***
factor(state)MN        2.111e-01  2.247e-02   9.395  < 2e-16 ***
factor(state)MO        1.050e-01  2.091e-02   5.022 5.42e-07 ***
factor(state)MS       -3.268e-02  2.711e-02  -1.205 0.228136
factor(state)MT        1.500e-01  2.207e-02   6.793 1.32e-11 ***
factor(state)NC        1.680e-01  2.360e-02   7.120 1.36e-12 ***
factor(state)ND        1.056e-01  2.205e-02   4.786 1.78e-06 ***
factor(state)NE        5.492e-02  2.099e-02   2.617 0.008926 **
factor(state)NH        3.920e-01  6.827e-02   5.742 1.03e-08 ***
factor(state)NJ        1.912e-01  4.196e-02   4.557 5.41e-06 ***
factor(state)NM       -2.018e-03  5.002e-02  -0.040 0.967814
factor(state)NV       -1.142e-02  5.172e-02  -0.221 0.825270
factor(state)NY        2.717e-01  2.377e-02  11.427  < 2e-16 ***
factor(state)OH        1.231e-01  2.192e-02   5.617 2.12e-08 ***
factor(state)OK        2.312e-02  2.952e-02   0.783 0.433595
factor(state)OR        2.248e-01  3.116e-02   7.214 6.90e-13 ***
factor(state)PA        1.629e-01  2.263e-02   7.198 7.75e-13 ***
factor(state)RI        3.878e-01  7.860e-02   4.934 8.52e-07 ***
factor(state)SC        7.413e-02  3.739e-02   1.982 0.047534 *
factor(state)SD        1.501e-01  2.160e-02   6.948 4.55e-12 ***
factor(state)TN        1.160e-01  2.094e-02   5.543 3.24e-08 ***
factor(state)TX       -8.246e-02  2.173e-02  -3.795 0.000151 ***
factor(state)UT        1.440e-01  5.686e-02   2.533 0.011377 *
factor(state)VA        1.219e-01  2.272e-02   5.364 8.77e-08 ***
factor(state)VT        4.690e-01  6.330e-02   7.410 1.64e-13 ***
factor(state)WA        3.214e-01  2.882e-02  11.153  < 2e-16 ***
factor(state)WI        2.801e-01  2.139e-02  13.095  < 2e-16 ***
factor(state)WV        1.253e-01  2.446e-02   5.123 3.21e-07 ***
factor(state)WY       -3.563e-02  3.934e-02  -0.906 0.365090
NonWhite               4.939e-03  6.065e-04   8.145 5.58e-16 ***
```

```
case_pct                 -8.977e-01  1.418e-01  -6.329 2.85e-10 ***
Black                     3.044e-03  4.521e-04   6.733 1.99e-11 ***
turnout_difference        1.620e-07  3.422e-08   4.735 2.30e-06 ***
Income                    2.634e-06  1.453e-07  18.125  < 2e-16 ***
Men_pct                  -8.661e-01  6.490e-02 -13.346  < 2e-16 ***
Transit                   9.703e-03  8.246e-04  11.767  < 2e-16 ***
factor(state)AR:NonWhite -2.859e-03  6.811e-04  -4.198 2.78e-05 ***
factor(state)AZ:NonWhite  1.477e-03  1.261e-03   1.171 0.241865
factor(state)CA:NonWhite -3.795e-03  7.995e-04  -4.747 2.17e-06 ***
factor(state)CO:NonWhite -1.494e-03  9.124e-04  -1.637 0.101738
factor(state)CT:NonWhite  2.934e-03  2.840e-03   1.033 0.301665
factor(state)DC:NonWhite         NA         NA      NA       NA
factor(state)DE:NonWhite  2.058e-03  6.560e-03   0.314 0.753737
factor(state)FL:NonWhite -1.361e-03  8.136e-04  -1.673 0.094443 .
factor(state)GA:NonWhite -3.642e-05  5.788e-04  -0.063 0.949830
factor(state)HI:NonWhite -1.595e-02  6.469e-03  -2.466 0.013721 *
factor(state)IA:NonWhite  1.509e-03  1.336e-03   1.130 0.258703
factor(state)ID:NonWhite -3.838e-03  1.227e-03  -3.128 0.001779 **
factor(state)IL:NonWhite  2.576e-04  8.867e-04   0.290 0.771464
factor(state)IN:NonWhite  2.788e-03  1.217e-03   2.292 0.021979 *
factor(state)KS:NonWhite -1.527e-03  8.307e-04  -1.839 0.066062 .
factor(state)KY:NonWhite  1.749e-03  1.373e-03   1.273 0.203002
factor(state)LA:NonWhite  1.117e-03  8.530e-04   1.310 0.190413
factor(state)MA:NonWhite -6.541e-03  3.232e-03  -2.024 0.043100 *
factor(state)MD:NonWhite -2.328e-03  1.018e-03  -2.286 0.022314 *
factor(state)ME:NonWhite  2.905e-02  1.367e-02   2.124 0.033740 *
factor(state)MI:NonWhite -5.177e-04  1.160e-03  -0.446 0.655541
factor(state)MN:NonWhite  3.195e-04  1.239e-03   0.258 0.796483
factor(state)MO:NonWhite  7.140e-04  1.215e-03   0.588 0.556678
factor(state)MS:NonWhite  1.292e-05  6.132e-04   0.021 0.983195
factor(state)MT:NonWhite  9.951e-04  8.989e-04   1.107 0.268381
factor(state)NC:NonWhite -2.317e-03  6.466e-04  -3.583 0.000345 ***
factor(state)ND:NonWhite  8.075e-04  8.813e-04   0.916 0.359568
factor(state)NE:NonWhite  1.088e-03  9.635e-04   1.129 0.258790
factor(state)NH:NonWhite -1.387e-02  9.639e-03  -1.439 0.150186
factor(state)NJ:NonWhite -2.921e-03  1.114e-03  -2.622 0.008787 **
factor(state)NM:NonWhite  1.734e-03  1.004e-03   1.727 0.084292 .
factor(state)NV:NonWhite  1.084e-03  1.795e-03   0.604 0.546041
factor(state)NY:NonWhite -7.948e-03  1.006e-03  -7.902 3.84e-15 ***
factor(state)OH:NonWhite  4.054e-03  1.258e-03   3.224 0.001278 **
factor(state)OK:NonWhite -5.149e-04  1.136e-03  -0.453 0.650405
factor(state)OR:NonWhite -3.386e-04  1.455e-03  -0.233 0.816016
factor(state)PA:NonWhite  3.089e-04  1.047e-03   0.295 0.767959
factor(state)RI:NonWhite  3.898e-03  3.672e-03   1.062 0.288522
factor(state)SC:NonWhite -9.033e-04  8.756e-04  -1.032 0.302323
factor(state)SD:NonWhite  1.584e-04  7.534e-04   0.210 0.833501
factor(state)TN:NonWhite -2.062e-03  8.232e-04  -2.504 0.012317 *
factor(state)TX:NonWhite  5.121e-04  6.611e-04   0.775 0.438616
factor(state)UT:NonWhite  3.223e-04  2.180e-03   0.148 0.882485
factor(state)VA:NonWhite -1.545e-03  7.028e-04  -2.198 0.027994 *
factor(state)VT:NonWhite -1.351e-02  1.410e-02  -0.958 0.337939
factor(state)WA:NonWhite -4.851e-03  1.080e-03  -4.493 7.31e-06 ***
factor(state)WI:NonWhite  3.008e-04  9.784e-04   0.307 0.758501
factor(state)WV:NonWhite  2.749e-03  3.007e-03   0.914 0.360642
factor(state)WY:NonWhite  7.235e-03  2.794e-03   2.590 0.009657 **
case_pct:Black            2.771e-02  7.101e-03   3.903 9.73e-05 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.0769 on 2915 degrees of freedom
Multiple R-squared:  0.7616,   Adjusted R-squared:  0.753
F-statistic: 88.71 on 105 and 2915 DF,  p-value: < 2.2e-16
```

Appendix 2: Correlation Matrix Heatmaps

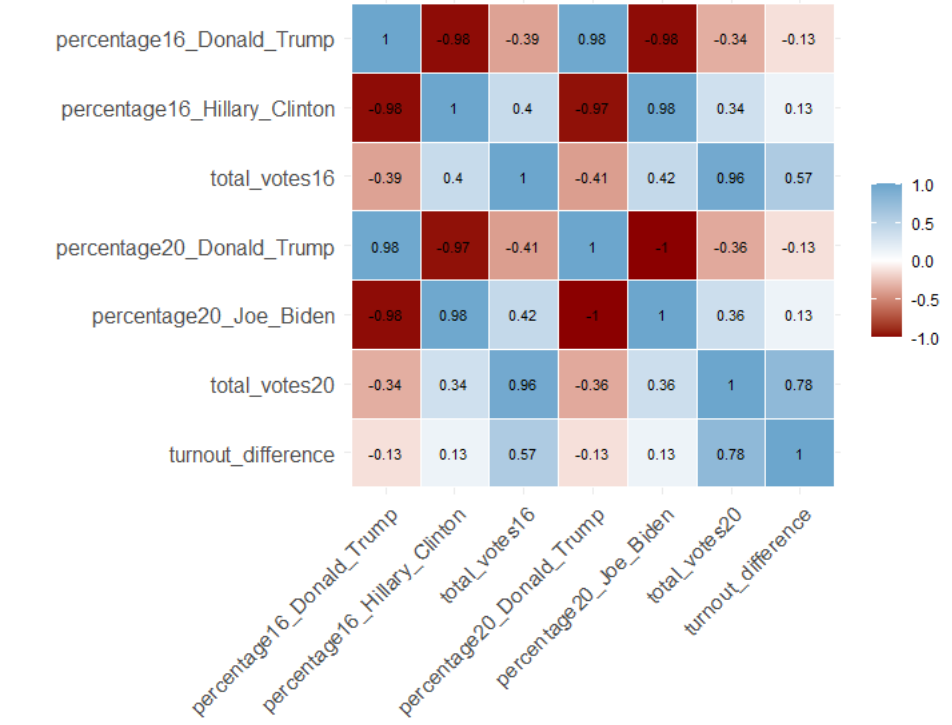

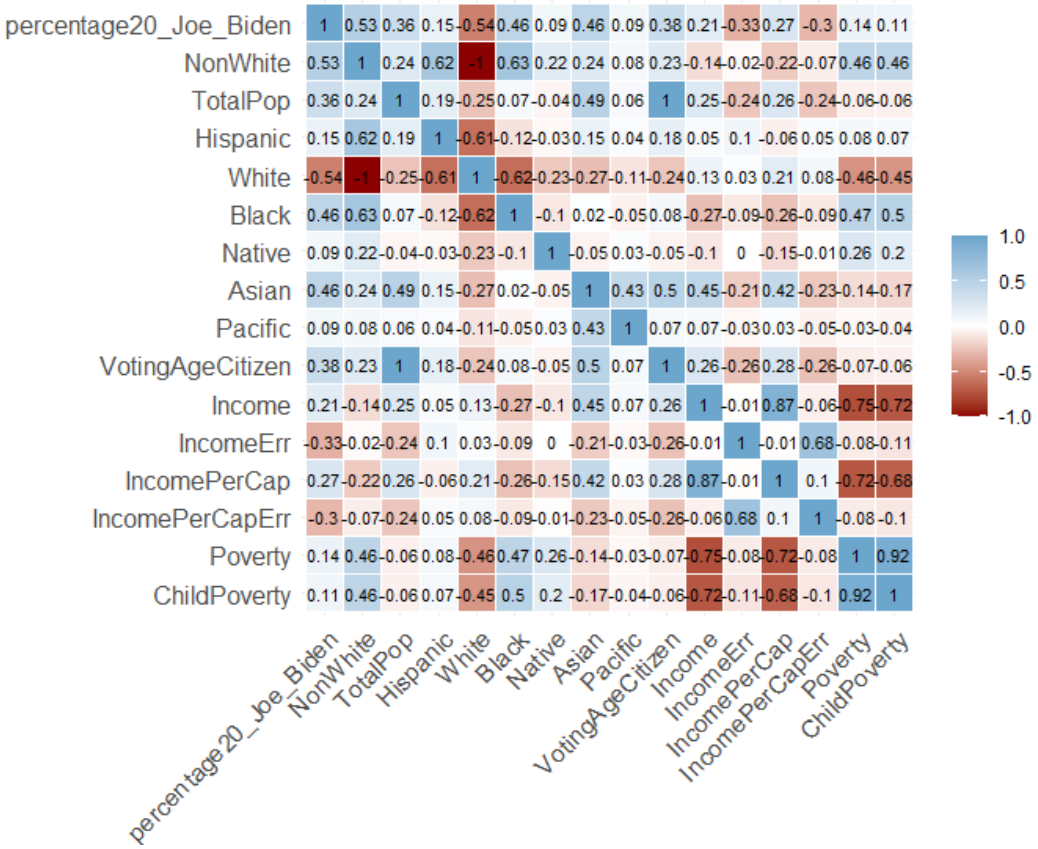

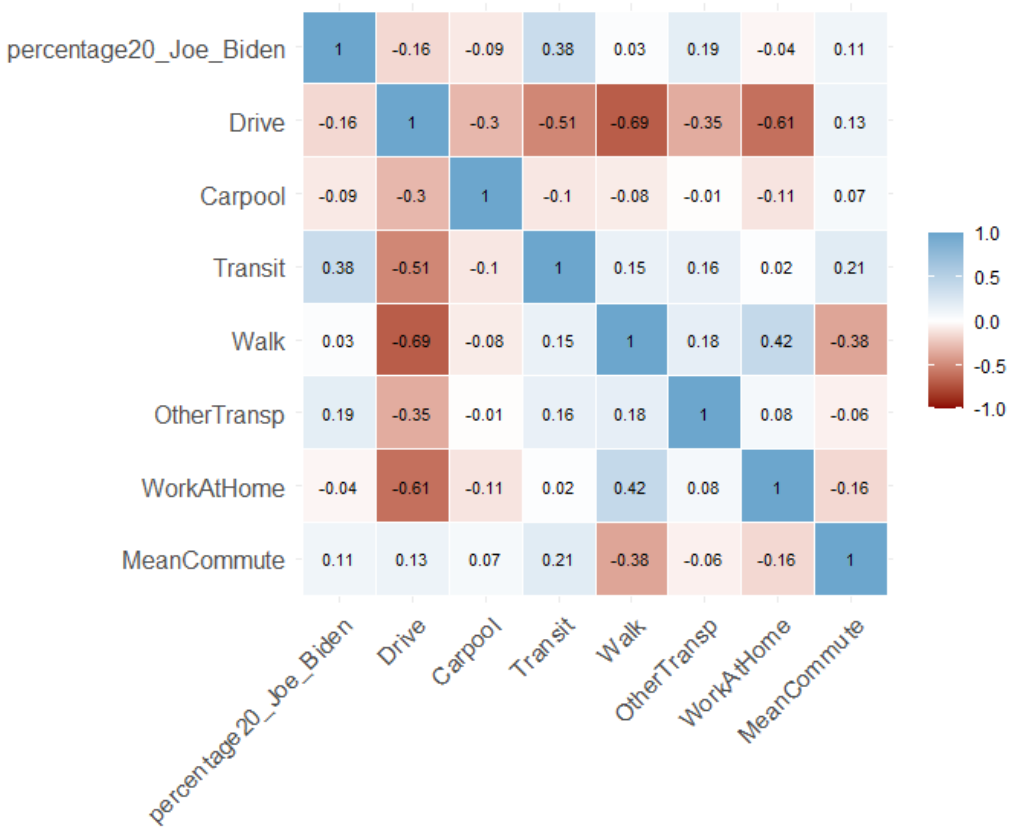
percentage20_Joe_Biden
Drive
Carpool
Transit
Walk
OtherTransp
WorkAtHome
MeanCommute
1 -0.16 -0.09 0.38 0.03 0.19 -0.04 0.11
-0.16 1 -0.3 -0.51 -0.69 -0.35 -0.61 0.13
-0.09 -0.3 1 -0.1 -0.08 -0.01 -0.11 0.07
0.38 -0.51 -0.1 1 0.15 0.16 0.02 0.21
0.03 -0.69 -0.08 0.15 1 0.18 0.42 -0.38
0.19 -0.35 -0.01 0.16 0.18 1 0.08 -0.06
-0.04 -0.61 -0.11 0.02 0.42 0.08 1 -0.16
0.11 0.13 0.07 0.21 -0.38 -0.06 -0.16 1
1.0
0.5
0.0
-0.5
-1.0

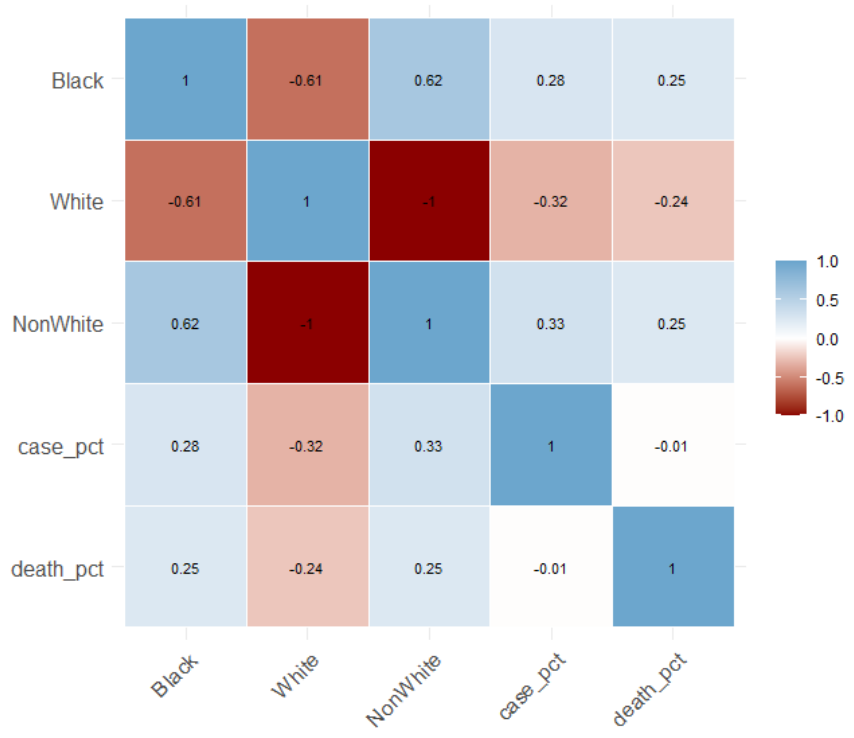
Black
White
NonWhite
case_pct
death_pct
1 -0.61 0.62 0.28 0.25
-0.61 1 -1 -0.32 -0.24
0.62 -1 1 0.33 0.25
0.28 -0.32 0.33 1 -0.01
0.25 -0.24 0.25 -0.01 1
1.0
0.5
0.0
-0.5
-1.0

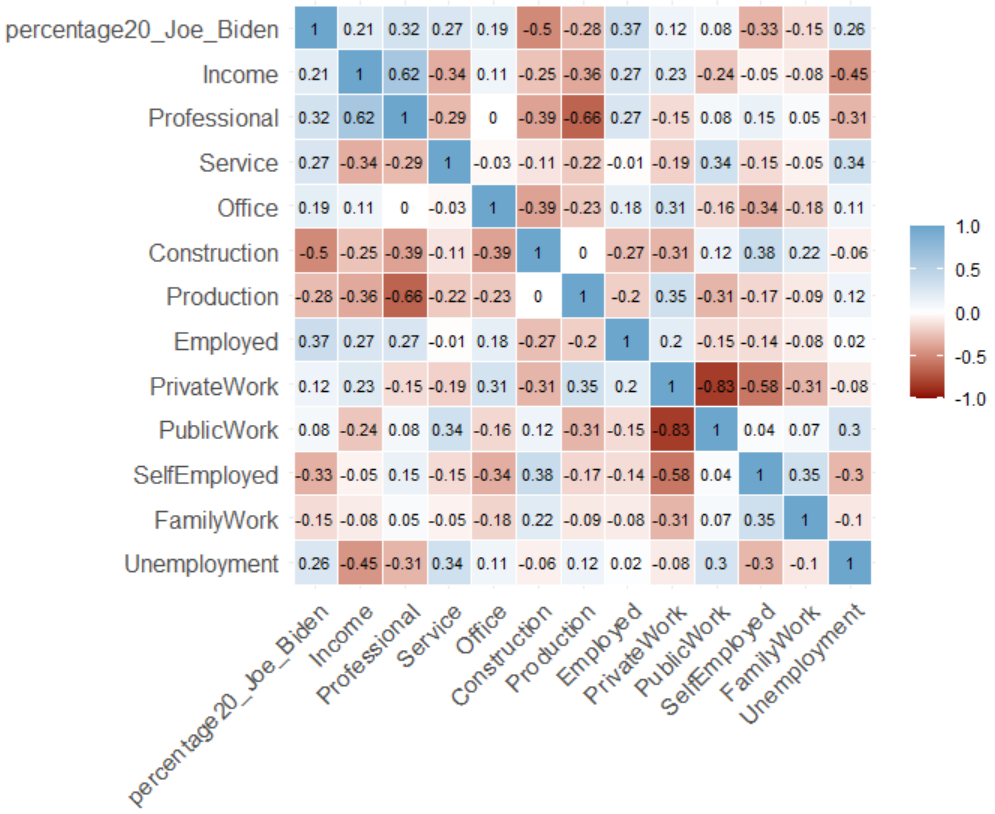

Appendix 3: Heteroscedasticity in Model

Heteroscedasticity looks present in multiple explanatory variables, including NonWhite, Black, turnout_difference, and Transit.

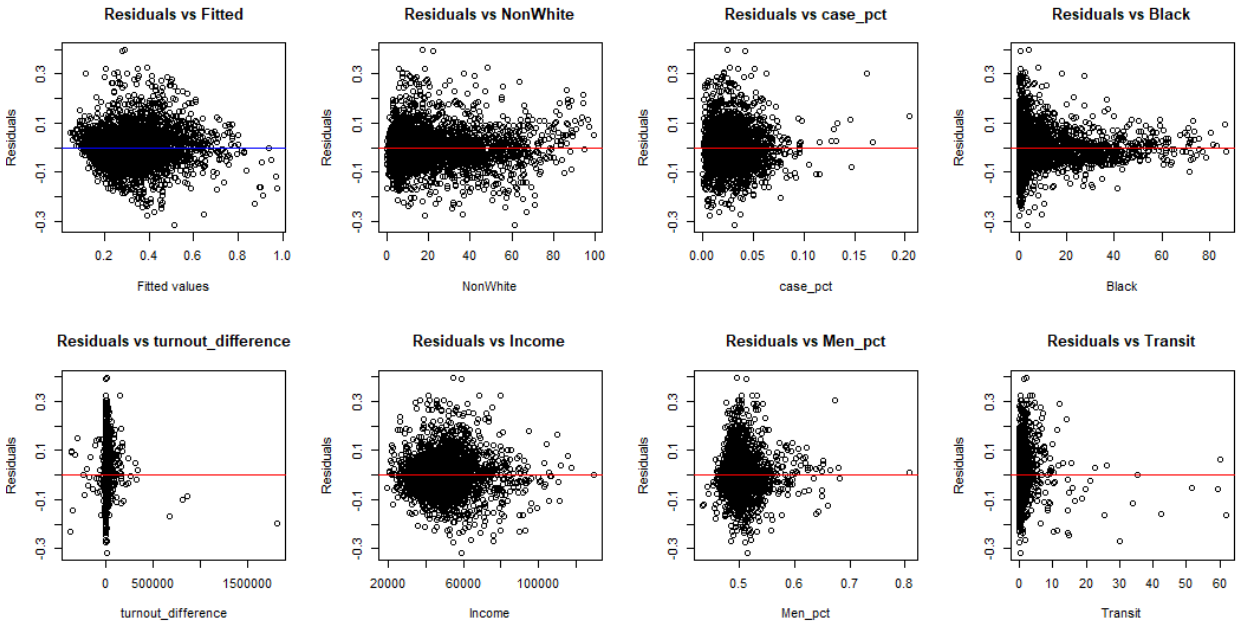

Appendix 4: Log Transformed Model

We take the log of the response variable to mitigate above heteroscedasticity. The new model has R^2 of 0.70 and the same 107 parameters.

log(percentage20_Joe_Biden) ~ factor(state)*NonWhite + case_pct*Black + turnout_difference + Income + Men_pct + Transit